

Практическое задание №1

Установка необходимых пакетов:

```
In [1]: !pip install -q tqdm
!pip install --upgrade --no-cache-dir gdown

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: gdown in /usr/local/lib/python3.8/dist-packages (4.4.0)
Collecting gdown
  Downloading gdown-4.5.4-py3-none-any.whl (14 kB)
Requirement already satisfied: six in /usr/local/lib/python3.8/dist-packages (from gdown) (1.15.0)
Requirement already satisfied: requests[socks] in /usr/local/lib/python3.8/dist-packages (from gdown) (2.23.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.8/dist-packages (from gdown) (4.64.1)
Requirement already satisfied: filelock in /usr/local/lib/python3.8/dist-packages (from gdown) (3.8.0)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.8/dist-packages (from gdown) (4.6.3)
Requirement already satisfied: urllib3!=1.25.0,!=1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.8/dist-packages (from requests[socks]->gdown) (1.24.3)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packages (from requests[socks]->gdown) (2.10)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.8/dist-packages (from requests[socks]->gdown) (3.0.4)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.8/dist-packages (from requests[socks]->gdown) (2022.9.24)
Requirement already satisfied: PySocks!=1.5.7,>=1.5.6 in /usr/local/lib/python3.8/dist-packages (from requests[socks]->gdown) (1.7.1)
Installing collected packages: gdown
  Attempting uninstall: gdown
    Found existing installation: gdown 4.4.0
    Uninstalling gdown-4.4.0:
      Successfully uninstalled gdown-4.4.0
  Successfully installed gdown-4.5.4
```

Монтирование Вашего Google Drive к текущему окружению:

```
In [2]: from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

Константы, которые пригодятся в коде далее, и ссылки (gdrive идентификаторы) на предоставляемые наборы данных:

```
In [3]: EVALUATE_ONLY = True
TEST_ON_LARGE_DATASET = True
TISSUE_CLASSES = ('ADI', 'BACK', 'DEB', 'LYM', 'MUC', 'MUS', 'NORM', 'STR', 'TUM')
DATASETS_LINKS = {
    'train': '1XtQzVQ5XbrfxpLHJuL0XBGJ5U7CS-cLi',
    'train_small': '1qd45xXfDwdZjktLFwQb-et-mAaFeCz0R',
    'train_tiny': '1I-2Z0uXLd4QwhZQ0ltp817Kn3J0Xgbui',
    'test': '1RfPou3pFKpuHDJZ-D9XDFzgvwpUBF1Dr',
    'test_small': '1wbRsog0n7uG1HIPGLhyN-PMET2kdQ21I',
    'test_tiny': '1viiB0s041CNsAK4itvX8PnYthJ-MDnQc'
}
```

Импорт необходимых зависимостей:

```
In [4]: from pathlib import Path
import numpy as np
from typing import List
#from tqdm.notebook import tqdm
#from time import sleep
from PIL import Image
import IPython.display
from sklearn.metrics import balanced_accuracy_score
import gdown
```

Класс Dataset

Предназначен для работы с наборами данных, обеспечивает чтение изображений и соответствующих меток, а также формирование пакетов (батчей).

```
In [5]: class Dataset:

    def __init__(self, name):
        self.name = name
        self.is_loaded = False
        url = f"https://drive.google.com/uc?export=download&confirm=pbef&id={DATASETS_LINKS[name]}"
        output = f'{name}.npz'
        gdown.download(url, output, quiet=False)
        print(f'Loading dataset {self.name} from npz.')
```

```
np_obj = np.load(f'{name}.npz')
self.images = np_obj['data']
self.labels = np_obj['labels']
self.n_files = self.images.shape[0]
self.is_loaded = True
print(f'Done. Dataset {name} consists of {self.n_files} images.')

def image(self, i):
    # read i-th image in dataset and return it as numpy array
    if self.is_loaded:
        return self.images[i, :, :, :]

def images_seq(self, n=None):
    # sequential access to images inside dataset (is needed for testing)
    for i in range(self.n_files if not n else n):
        yield self.image(i)

def random_image_with_label(self):
    # get random image with label from dataset
    i = np.random.randint(self.n_files)
    return self.image(i), self.labels[i]

def random_batch_with_labels(self, n):
    # create random batch of images with labels (is needed for training)
    indices = np.random.choice(self.n_files, n)
    imgs = []
    for i in indices:
        img = self.image(i)
        imgs.append(self.image(i))
    logits = np.array([self.labels[i] for i in indices])
    return np.stack(imgs), logits

def image_with_label(self, i: int):
    # return i-th image with label from dataset
    return self.image(i), self.labels[i]
```

Пример использования класса Dataset

Загрузим обучающий набор данных, получим произвольное изображение с меткой. После чего визуализируем изображение, выведем метку. В будущем, этот кусок кода можно закомментировать или убрать.

```
In [6]: d_train_tiny = Dataset('train_tiny')

img, lbl = d_train_tiny.random_image_with_label()
print()
print(f'Got numpy array of shape {img.shape}, and label with code {lbl}.')
print(f'Label code corresponds to {TISSUE_CLASSES[lbl]} class.')

pil_img = Image.fromarray(img)
IPython.display.display(pil_img)
```

Downloading...

From: <https://drive.google.com/uc?export=download&confirm=pbef&id=1I-2Z0uXLd4QwhZQqltp817Kn3J0Xgbui>

To: /content/train_tiny.npz


100%|██████████| 105M/105M [00:00<00:00, 122MB/s]

Loading dataset train_tiny from npz.

Done. Dataset train_tiny consists of 900 images.

Got numpy array of shape (224, 224, 3), and label with code 0.

Label code corresponds to ADI class.



Класс Metrics

Реализует метрики точности, используемые для оценивания модели:

1. точность,
2. сбалансированную точность.

```
In [7]: class Metrics:

    @staticmethod
    def accuracy(gt: List[int], pred: List[int]):
        assert len(gt) == len(pred), 'gt and prediction should be of equal length'
        return sum(int(i[0] == i[1]) for i in zip(gt, pred)) / len(gt)

    @staticmethod
    def accuracy_balanced(gt: List[int], pred: List[int]):
        return balanced_accuracy_score(gt, pred)

    @staticmethod
    def print_all(gt: List[int], pred: List[int], info: str):
        print(f'metrics for {info}:')
        print('\t accuracy {:.4f}:'.format(Metrics.accuracy(gt, pred)))
        print('\t balanced accuracy {:.4f}:'.format(Metrics.accuracy_balanced(gt, pred)))
```

Класс Model

Класс, хранящий в себе всю информацию о модели.

Вам необходимо реализовать методы `save`, `load` для сохранения и загрузки модели. Особенно актуально это будет во время тестирования на дополнительных наборах данных.

Пожалуйста, убедитесь, что сохранение и загрузка модели работает корректно. Для этого обучите модель, протестируйте, сохраните ее в файл, перезапустите среду выполнения, загрузите обученную модель из файла, вновь протестируйте ее на тестовой выборке и убедитесь в том, что получаемые метрики совпадают с полученными для тестовой выборки ранее.

Также, Вы можете реализовать дополнительные функции, такие как:

1. валидацию модели на части обучающей выборки;
2. использование кроссвалидации;
3. автоматическое сохранение модели при обучении;
4. загрузку модели с какой-то конкретной итерации обучения (если используется итеративное обучение);
5. вывод различных показателей в процессе обучения (например, значение функции потерь на каждой эпохе);

6. построение графиков, визуализирующих процесс обучения (например, график зависимости функции потерь от номера эпохи обучения);
7. автоматическое тестирование на тестовом наборе/наборах данных после каждой эпохи обучения (при использовании итеративного обучения);
8. автоматический выбор гиперпараметров модели во время обучения;
9. сохранение и визуализацию результатов тестирования;
10. Использование аугментации и других способов синтетического расширения набора данных (дополнительным плюсом будет обоснование необходимости и обоснование выбора конкретных типов аугментации)
11. и т.д.

Полный список опций и дополнений приведен в презентации с описанием задания.

При реализации дополнительных функций допускается добавление параметров в существующие методы и добавление новых методов в класс модели.

```
In [8]: import keras
import tensorflow as tf
from keras import layers
from keras.applications import EfficientNetV2S
from sklearn.model_selection import train_test_split
from keras import callbacks
import matplotlib.pyplot as plt
```

```
In [9]: from tensorflow.python.util.tf_export import kwarg_only
class Model:

    def __init__(self):
        # аугментация
        data_augmentation = keras.Sequential(
            [
                layers.RandomFlip(),
                layers.RandomRotation(0.2),
                layers.RandomZoom((-0.2, 0)),
            ]
        )

        self.model_1_epoch_num = 40

        self.model_1 = keras.models.Sequential()
```

```
self.model_1.add(layers.Input(shape=(224, 224, 3)))

# augmentation layer
self.model_1.add(data_augmentation)

self.model_1.add(layers.Normalization())
self.model_1.add(EfficientNetV2S(include_top=False, input_shape=(224, 224, 3)))
self.model_1.add(layers.Flatten())
self.model_1.add(layers.Dense(9, activation='softmax'))

self.model_1.compile(optimizer=keras.optimizers.Adam(),
                    loss=keras.losses.SparseCategoricalCrossentropy(from_logits=False),
                    metrics=['sparse_categorical_accuracy'])

def save(self, name: str):
    print(f'saving started')
    self.model_1.save(f'/content/drive/MyDrive/{name}.hdf5')
    print(f'saving done')

def load(self, name: str = 'best'):
    print(f'loading started')
    name_to_id_dict = {
        'best': '1-8StHduxeVrtZAE1-5AQIPT0YQj11Rso'
    }
    output = f'{name}.hdf5'
    gdown.download(f'https://drive.google.com/uc?id={name_to_id_dict[name]}', output, quiet=False)
    self.model_1 = tf.keras.models.load_model(f'{name}.hdf5')
    print(f'loading done')

def loading_history(self, name: str = 'model_1_cur_logs'):
    print(f'loading started')
    name_to_id_dict = {
        'model_1_cur_logs': '1-49Mm_4ePquWz9ydjz6V54N807xVMc7X'
    }
    output = f'{name}.csv'
    gdown.download(f'https://drive.google.com/uc?id={name_to_id_dict[name]}', output, quiet=False)
    self.history_1 = pd.read_csv(f'{name}.csv')
    print(f'loading done')

def train(self, dataset: Dataset, val_dataset = None):
    # you can add some plots for better visualization,
    # you can add model autosaving during training,
    # etc.
    print(f'training started')
```

```

if val_dataset is None:
    self.train_1(dataset.images, dataset.labels)
else:
    self.train_1(dataset.images, dataset.labels, (val_dataset.images, val_dataset.labels))
print(f'training done')

def train_1(self, images_train, labels_train, validation_data=None):

    # weights of different classes

    val, counts = np.unique(labels_train, return_counts=True)
    sort = np.argsort(val)
    val = val[sort]
    counts = 1 / (counts[sort] + 0.001) * labels_train.shape[0] / 2

    weights = dict(zip(val, counts))

    print("weights done")

    # splitting data

    if validation_data:
        images_train = images_train
        labels_train = labels_train
    else:
        # валидация
        images_train, images_val, labels_train, labels_val = train_test_split(images_train, labels_train,
                                                                                test_size=0.05, random_state=42)
        validation_data = (images_val, labels_val)

    # callbacks

    initial_learning_rate = 0.0001
    def lr_exp_decay(epoch, lr):
        k = 0.1
        return initial_learning_rate * np.exp(-k*epoch)

    # сохранение модели
    checkpoint = keras.callbacks.ModelCheckpoint('/content/drive/MyDrive/model_1_cur.hdf5', monitor='val_loss', save_freq='epoch',
                                                verbose=1, save_best_only=True, save_weights_only=False, mode='auto')

    # выбор гиперпараметров
    # (количество эпох)
    earlystopping = keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0.1, patience=10, verbose=1,
                                                mode='auto', baseline=None)

```



```
# (learning rate)
lrscheduler = keras.callbacks.LearningRateScheduler(lr_exp_decay, verbose=1)
reduce_lr = keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=5, verbose=1,
                                              mode='auto', min_delta=0.01, cooldown=0, min_lr=0)

# сохранение показателей
csvlogger = keras.callbacks.CSVLogger('/content/drive/MyDrive/model_1_cur_logs.csv', separator=',', append=False)

callbacks = [lrscheduler, checkpoint, earlystopping, reduce_lr, csvlogger]

# train
# вывод показателей
# (fit выводит показатели для каждой эпохи)
self.train = self.model_1.fit(images_train, labels_train,
                              epochs=self.model_1_epoch_num,
                              class_weight=weights,
                              validation_data=validation_data,
                              callbacks=callbacks)

self.history_1 = self.train.history
self.model_1.save(f'/content/drive/MyDrive/cur_trial_1.hdf5')

print("model_1 done")

def test_on_dataset(self, dataset: Dataset, limit=None):
    # you can upgrade this code if you want to speed up testing using batches
    n = dataset.n_files if not limit else int(dataset.n_files * limit)
    predictions = self.test_on_dataset_1(dataset.images[:n])
    return predictions

def test_on_dataset_1(self, images):
    predictions_1 = tf.argmax(self.model_1.predict(images), axis=1)
    return predictions_1

def test_on_image(self, img: np.ndarray):
    # todo: replace this code
    img_array = tf.keras.utils.img_to_array(img)
    img_array = tf.expand_dims(img_array, 0)
    prediction = self.model.predict(img_array)
    return prediction

def learning_plots(self):
    acc = self.history_1['sparse_categorical_accuracy']
    val_acc = self.history_1['val_sparse_categorical_accuracy']
```

```

loss = self.history_1['loss']
val_loss = self.history_1['val_loss']

lr = self.history_1['lr']

epochs_range = range(len(lr))

plt.figure(figsize=(15, 5))
plt.subplot(1, 3, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')
plt.grid()

plt.subplot(1, 3, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.grid()

plt.subplot(1, 3, 3)
plt.plot(epochs_range, lr, label='Learning rate')
plt.legend(loc='upper right')
plt.title('Learning rate')
plt.grid()
plt.show()

```

```

In [ ]: from sklearn.model_selection import KFold
        from sklearn.metrics import balanced_accuracy_score
        from collections import defaultdict

```

```

In [ ]: # кросс-валидация
        def cross_val(dataset: Dataset, n_splits=3):
            X = dataset.images
            y = dataset.labels

            scorer = balanced_accuracy_score
            cv = KFold(n_splits=n_splits, shuffle=True, random_state=42)
            res=[]

            t = cv.split(X)

```

```
for train_i, test_i in t:
    X_train = X[train_i]
    y_train = y[train_i]
    X_test = X[test_i]
    y_test = y[test_i]
    model = Model()
    model.train_1(images_train=X_train, labels_train=y_train, validation_data=(X_test, y_test))
    pred = model.test_on_dataset_1(X_test)
    res.append(scorer(y_test, pred))

splits_num = range(len(res))

plt.figure(figsize=(4, 3))
plt.plot(splits_num, res, label='accuracy')
plt.legend(loc='lower right')
plt.title('Cross validation accuracy')
plt.ylim([0,1])
plt.grid()
plt.show()

return res
```

Классификация изображений

Классификация изображений. Используем полные датасеты `train`, `test`.

```
In [ ]: d_train = Dataset('train')
```

```
Downloading...
From: https://drive.google.com/uc?export=download&confirm=pbef&id=18bNP8R_k8FYq0Gx0DnC2j_UE0aGqV114
To: /content/train.npz
100%|██████████| 2.10G/2.10G [00:09<00:00, 218MB/s]
Loading dataset train from npz.
Done. Dataset train consists of 18000 images.
```

Обучение модели

Используется аугментация данных стандартными функциями библиотеки `keras.layers`. Далее используется архитектура `EfficientNetV2S` с уже обученными на `ImageNet` слоями. После нее добавлен обычный полносвязный слой с 9 выходами и функцией активации `softmax` для окончательной классификации.

Модель обучена на 14 эпохах.

Использованы `callback`-функции:

- `LearningRateScheduler` - `learning rate` зависит от эпохи (экспоненциально)
- `ModelCheckpoint` - периодически сохраняет модель (только с лучшими весами)
- `EarlyStopping` - ранняя остановка для недопуска переобучения
- `ReduceLROnPlateau` - уменьшить `learning rate`, если лосс выходит на плато
- `CSVLogger` - сохранение информации об обучении (`loss`, `accuracy`, ...) в формате `.csv`

```
In [ ]: model = Model()
if not EVALUATE_ONLY:
    model.train(d_train)
    model.save('best')
else:
    model.load('best')
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/efficientnet_v2/efficientnetv2-s_notop.h5
82420632/82420632 [=====] - 5s 0us/step
training started
weights done

Epoch 1: LearningRateScheduler setting learning rate to 0.0001.

Epoch 1/40

535/535 [=====] - ETA: 0s - loss: 1.4519 - sparse_categorical_accuracy: 0.8936

Epoch 1: val_loss improved from inf to 0.15565, saving model to /content/drive/MyDrive/model_1_cur.hdf5

535/535 [=====] - 296s 497ms/step - loss: 1.4519 - sparse_categorical_accuracy: 0.8936 - val_loss: 0.1556 - val_sparse_categorical_accuracy: 0.9544 - lr: 1.0000e-04

Epoch 2: LearningRateScheduler setting learning rate to 9.048374180359596e-05.

Epoch 2/40

535/535 [=====] - ETA: 0s - loss: 0.5369 - sparse_categorical_accuracy: 0.9618

Epoch 2: val_loss improved from 0.15565 to 0.10841, saving model to /content/drive/MyDrive/model_1_cur.hdf5

535/535 [=====] - 263s 491ms/step - loss: 0.5369 - sparse_categorical_accuracy: 0.9618 - val_loss: 0.1084 - val_sparse_categorical_accuracy: 0.9633 - lr: 9.0484e-05

Epoch 3: LearningRateScheduler setting learning rate to 8.187307530779819e-05.

Epoch 3/40

535/535 [=====] - ETA: 0s - loss: 0.3666 - sparse_categorical_accuracy: 0.9727

Epoch 3: val_loss improved from 0.10841 to 0.06697, saving model to /content/drive/MyDrive/model_1_cur.hdf5

535/535 [=====] - 262s 490ms/step - loss: 0.3666 - sparse_categorical_accuracy: 0.9727 - val_loss: 0.0670 - val_sparse_categorical_accuracy: 0.9767 - lr: 8.1873e-05

Epoch 4: LearningRateScheduler setting learning rate to 7.408182206817179e-05.

Epoch 4/40

535/535 [=====] - ETA: 0s - loss: 0.2882 - sparse_categorical_accuracy: 0.9810

Epoch 4: val_loss improved from 0.06697 to 0.04043, saving model to /content/drive/MyDrive/model_1_cur.hdf5

535/535 [=====] - 264s 494ms/step - loss: 0.2882 - sparse_categorical_accuracy: 0.9810 - val_loss: 0.0404 - val_sparse_categorical_accuracy: 0.9856 - lr: 7.4082e-05

Epoch 5: LearningRateScheduler setting learning rate to 6.703200460356394e-05.

Epoch 5/40

535/535 [=====] - ETA: 0s - loss: 0.2172 - sparse_categorical_accuracy: 0.9851

Epoch 5: val_loss did not improve from 0.04043

535/535 [=====] - 258s 482ms/step - loss: 0.2172 - sparse_categorical_accuracy: 0.9851 - val_loss: 0.0734 - val_sparse_categorical_accuracy: 0.9733 - lr: 6.7032e-05

Epoch 6: LearningRateScheduler setting learning rate to 6.065306597126335e-05.

Epoch 6/40

535/535 [=====] - ETA: 0s - loss: 0.1813 - sparse_categorical_accuracy: 0.9874

Epoch 6: val_loss did not improve from 0.04043

535/535 [=====] - 258s 482ms/step - loss: 0.1813 - sparse_categorical_accuracy: 0.9874 - val_loss: 0.0729 - val_sparse_categorical_accuracy: 0.9767 - lr: 6.0653e-05

Epoch 7: LearningRateScheduler setting learning rate to 5.488116360940264e-05.

Epoch 7/40

535/535 [=====] - ETA: 0s - loss: 0.1367 - sparse_categorical_accuracy: 0.9905

Epoch 7: val_loss did not improve from 0.04043

535/535 [=====] - 258s 482ms/step - loss: 0.1367 - sparse_categorical_accuracy: 0.9905 - val_loss: 0.0883 - val_sparse_categorical_accuracy: 0.9756 - lr: 5.4881e-05

Epoch 8: LearningRateScheduler setting learning rate to 4.965853037914095e-05.

Epoch 8/40

535/535 [=====] - ETA: 0s - loss: 0.1493 - sparse_categorical_accuracy: 0.9899

Epoch 8: val_loss did not improve from 0.04043

535/535 [=====] - 257s 481ms/step - loss: 0.1493 - sparse_categorical_accuracy: 0.9899 - val_loss: 0.0738 - val_sparse_categorical_accuracy: 0.9822 - lr: 4.9659e-05

Epoch 9: LearningRateScheduler setting learning rate to 4.493289641172216e-05.

Epoch 9/40

535/535 [=====] - ETA: 0s - loss: 0.1060 - sparse_categorical_accuracy: 0.9921

Epoch 9: val_loss improved from 0.04043 to 0.03137, saving model to /content/drive/MyDrive/model_1_cur.hdf5

Epoch 9: ReduceLROnPlateau reducing learning rate to 2.246644908154849e-05.

535/535 [=====] - 263s 492ms/step - loss: 0.1060 - sparse_categorical_accuracy: 0.9921 - val_loss: 0.0314 - val_sparse_categorical_accuracy: 0.9878 - lr: 4.4933e-05

Epoch 10: LearningRateScheduler setting learning rate to 4.0656965974059915e-05.

Epoch 10/40

535/535 [=====] - ETA: 0s - loss: 0.0990 - sparse_categorical_accuracy: 0.9928

Epoch 10: val_loss did not improve from 0.03137

535/535 [=====] - 259s 485ms/step - loss: 0.0990 - sparse_categorical_accuracy: 0.9928 - val_loss: 0.0464 - val_sparse_categorical_accuracy: 0.9856 - lr: 4.0657e-05

Epoch 11: LearningRateScheduler setting learning rate to 3.678794411714424e-05.

Epoch 11/40

535/535 [=====] - ETA: 0s - loss: 0.0922 - sparse_categorical_accuracy: 0.9932

Epoch 11: val_loss did not improve from 0.03137

535/535 [=====] - 257s 481ms/step - loss: 0.0922 - sparse_categorical_accuracy: 0.9932 - val_loss: 0.0368 - val_sparse_categorical_accuracy: 0.9911 - lr: 3.6788e-05

Epoch 12: LearningRateScheduler setting learning rate to 3.3287108369807955e-05.

Epoch 12/40

535/535 [=====] - ETA: 0s - loss: 0.0587 - sparse_categorical_accuracy: 0.9953

Epoch 12: val_loss did not improve from 0.03137

```
535/535 [=====] - 258s 481ms/step - loss: 0.0587 - sparse_categorical_accuracy: 0.9953 - val_loss: 0.0500 - val_sparse_categorical_accuracy: 0.9889 - lr: 3.3287e-05
```

Epoch 13: LearningRateScheduler setting learning rate to 3.0119421191220204e-05.

Epoch 13/40

```
535/535 [=====] - ETA: 0s - loss: 0.0620 - sparse_categorical_accuracy: 0.9954
```

Epoch 13: val_loss did not improve from 0.03137

```
535/535 [=====] - 259s 485ms/step - loss: 0.0620 - sparse_categorical_accuracy: 0.9954 - val_loss: 0.0323 - val_sparse_categorical_accuracy: 0.9878 - lr: 3.0119e-05
```

Epoch 14: LearningRateScheduler setting learning rate to 2.725317930340126e-05.

Epoch 14/40

```
535/535 [=====] - ETA: 0s - loss: 0.0576 - sparse_categorical_accuracy: 0.9958
```

Epoch 14: val_loss did not improve from 0.03137

Epoch 14: ReduceLROnPlateau reducing learning rate to 1.3626589861814864e-05.

```
535/535 [=====] - 260s 485ms/step - loss: 0.0576 - sparse_categorical_accuracy: 0.9958 - val_loss: 0.0426 - val_sparse_categorical_accuracy: 0.9833 - lr: 2.7253e-05
```

Epoch 14: early stopping

model_1 done

training done

saving started

saving done

In []:

Тестируем на датасете test

In []: `d_test = Dataset('test')`

Downloading...

From: <https://drive.google.com/uc?export=download&confirm=pbef&id=14lLsHGRxcTlMkX2Y5j9KxH1W62ZjSiye>

To: /content/test.npz

100%|██████████| 525M/525M [00:11<00:00, 44.6MB/s]

Loading dataset test from npz.

Done. Dataset test consists of 4500 images.

In []: `pred_1 = model.test_on_dataset(d_test, limit=1)`
`Metrics.print_all(d_test.labels[:len(pred_1)], pred_1, 'full test')`

```
141/141 [=====] - 20s 124ms/step
metrics for full test:
    accuracy 0.9916:
    balanced accuracy 0.9916:
```

Чтение данных

Есть возможность считывать веса и данные обучения лучшей модели.

```
In [ ]: model_1 = Model()
        model_1.load()

loading started
Downloading...
From: https://drive.google.com/uc?id=1-8StHduxeVrtZAE1-5AQIPT0YQj11Rso
To: /content/best.hdf5
100%|██████████| 251M/251M [00:01<00:00, 152MB/s]
loading done
```

```
In [ ]: pred_1 = model_1.test_on_dataset(d_test, limit=1)
        Metrics.print_all(d_test.labels[:len(pred_1)], pred_1, 'full test')

141/141 [=====] - 20s 125ms/step
metrics for full test:
    accuracy 0.9916:
    balanced accuracy 0.9916:
```

```
In [ ]: import pandas as pd
```

```
In [ ]: model_1.loading_history()

loading started
Downloading...
From: https://drive.google.com/uc?id=1-49Mm_4ePquWz9ydz6V54N807xVMc7X
To: /content/model_1_cur_logs.csv
100%|██████████| 1.40k/1.40k [00:00<00:00, 2.91MB/s]
loading done
```

```
In [ ]: model_1.history_1.head()
```


Out[]:

| | epoch | loss | lr | sparse_categorical_accuracy | val_loss | val_sparse_categorical_accuracy |
|----------|-------|----------|----------|-----------------------------|----------|---------------------------------|
| 0 | 0 | 1.451855 | 0.000100 | 0.893626 | 0.155648 | 0.954444 |
| 1 | 1 | 0.536886 | 0.000090 | 0.961813 | 0.108412 | 0.963333 |
| 2 | 2 | 0.366590 | 0.000082 | 0.972749 | 0.066970 | 0.976667 |
| 3 | 3 | 0.288158 | 0.000074 | 0.980994 | 0.040435 | 0.985556 |
| 4 | 4 | 0.217177 | 0.000067 | 0.985146 | 0.073438 | 0.973333 |

Визуализации

Матрица ошибок для выборки `test`.

In [10]:

```
final_model = Model()
final_model.load('best')
d_test = Dataset('test')
pred_1 = final_model.test_on_dataset(d_test)
Metrics.print_all(d_test.labels, pred_1, 'test')
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/efficientnet_v2/efficientnetv2-s_notop.h5
82420632/82420632 [=====] - 1s 0us/step
loading started


Downloading...
From: <https://drive.google.com/uc?id=1-8StHduxeVrtZAE1-5AQIPT0YQj11Rso>
To: /content/best.hdf5
100%|██████████| 251M/251M [00:01<00:00, 127MB/s]
loading done

Downloading...
From: <https://drive.google.com/uc?export=download&confirm=pbef&id=1RfPou3pFKpuHDJZ-D9XDFzgvwpUBF1Dr>
To: /content/test.npz
100%|██████████| 525M/525M [00:04<00:00, 107MB/s]

Loading dataset test from npz.
 Done. Dataset test consists of 4500 images.
 141/141 [=====] - 720s 5s/step
 metrics for test:
 accuracy 0.9916:
 balanced accuracy 0.9916:

```
In [11]: # матрица ошибок
import sklearn
conf_matr = sklearn.metrics.confusion_matrix(d_test.labels[:len(pred_1)], pred_1)
disp = sklearn.metrics.ConfusionMatrixDisplay(conf_matr)
disp.plot()
```

Out[11]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7feae57271c0>



```
In [ ]: mcm = sklearn.metrics.multilabel_confusion_matrix(d_test.labels[:len(pred_1)], pred_1)
tn = mcm[:, 0, 0]
tp = mcm[:, 1, 1]
fn = mcm[:, 1, 0]
fp = mcm[:, 0, 1]
print("Sensitivity for each label:")
print(tp / (tp + fn))
print("Specificity for each label:")
print(tn / (tn + fp))
```

Sensitivity for each label:
 [0.992 0.996 0.986 1. 0.998 0.996 0.988 0.972 0.996]
 Specificity for each label:
 [1. 0.9995 0.9995 1. 0.998 0.997 0.99925 0.9985 0.99875]

Параметры (`loss` , `accuracy` , `learning rate`) в процессе обучения

```
In [ ]: final_model.loading_history()
```

loading started

Downloading...




From: https://drive.google.com/uc?id=1-49Mm_4ePquWz9ydz6V54N807xVMc7X

To: /content/model_1_cur_logs.csv

100%|██████████| 1.40k/1.40k [00:00<00:00, 3.68MB/s]

loading done

```
In [ ]: # графики обучения
final_model.learning_plots()
```



```
In [ ]:
```

Кросс-валидация

Написана также функция кросс-валидации для модели.

Ниже приведен пример работы функции. На полном датасете проводить кросс-валидацию слишком долго, и, в общем-то, не особо имеет смысл.

Модель обучается на 5 эпохах на датасете train_tiny.

```
In [ ]: d_train_cross = Dataset('train_tiny')
```

```
Downloading...  
From: https://drive.google.com/uc?export=download&confirm=pbef&id=1zdQ4BoKXR-bRU0j0_ja5LeDjl_bPtbg0  
To: /content/train_tiny.npz  
100%|██████████| 105M/105M [00:00<00:00, 112MB/s]  
Loading dataset train_tiny from npz.  
Done. Dataset train_tiny consists of 900 images.
```

```
In [ ]: cross_val_scores = cross_val(dataset=d_train_cross, n_splits=5)
```

weights done

Epoch 1: LearningRateScheduler setting learning rate to 0.001.

Epoch 1/5

23/23 [=====] - ETA: 0s - loss: 6.3189 - sparse_categorical_accuracy: 0.6583

Epoch 1: val_loss improved from inf to 1.20344, saving model to /content/drive/MyDrive/model_1_cur.hdf5

23/23 [=====] - 41s 907ms/step - loss: 6.3189 - sparse_categorical_accuracy: 0.6583 - val_loss: 1.2034
- val_sparse_categorical_accuracy: 0.8111 - lr: 0.0010

Epoch 2: LearningRateScheduler setting learning rate to 0.0009048374180359595.

Epoch 2/5

23/23 [=====] - ETA: 0s - loss: 4.0421 - sparse_categorical_accuracy: 0.8139

Epoch 2: val_loss improved from 1.20344 to 0.46617, saving model to /content/drive/MyDrive/model_1_cur.hdf5

23/23 [=====] - 16s 696ms/step - loss: 4.0421 - sparse_categorical_accuracy: 0.8139 - val_loss: 0.4662
- val_sparse_categorical_accuracy: 0.8889 - lr: 9.0484e-04

Epoch 3: LearningRateScheduler setting learning rate to 0.0008187307530779819.

Epoch 3/5

23/23 [=====] - ETA: 0s - loss: 2.7973 - sparse_categorical_accuracy: 0.8597

Epoch 3: val_loss improved from 0.46617 to 0.42818, saving model to /content/drive/MyDrive/model_1_cur.hdf5

23/23 [=====] - 16s 690ms/step - loss: 2.7973 - sparse_categorical_accuracy: 0.8597 - val_loss: 0.4282
- val_sparse_categorical_accuracy: 0.9056 - lr: 8.1873e-04

Epoch 4: LearningRateScheduler setting learning rate to 0.0007408182206817179.

Epoch 4/5

23/23 [=====] - ETA: 0s - loss: 2.2126 - sparse_categorical_accuracy: 0.8972

Epoch 4: val_loss did not improve from 0.42818

23/23 [=====] - 12s 519ms/step - loss: 2.2126 - sparse_categorical_accuracy: 0.8972 - val_loss: 1.0566
- val_sparse_categorical_accuracy: 0.8222 - lr: 7.4082e-04

Epoch 5: LearningRateScheduler setting learning rate to 0.0006703200460356394.

Epoch 5/5

23/23 [=====] - ETA: 0s - loss: 2.3558 - sparse_categorical_accuracy: 0.9000

Epoch 5: val_loss did not improve from 0.42818

23/23 [=====] - 12s 517ms/step - loss: 2.3558 - sparse_categorical_accuracy: 0.9000 - val_loss: 0.5252
- val_sparse_categorical_accuracy: 0.9056 - lr: 6.7032e-04

model_1 done

6/6 [=====] - 3s 129ms/step

weights done

Epoch 1: LearningRateScheduler setting learning rate to 0.001.

Epoch 1/5

23/23 [=====] - ETA: 0s - loss: 7.4492 - sparse_categorical_accuracy: 0.6431

Epoch 1: val_loss improved from inf to 1.53536, saving model to /content/drive/MyDrive/model_1_cur.hdf5

23/23 [=====] - 43s 853ms/step - loss: 7.4492 - sparse_categorical_accuracy: 0.6431 - val_loss: 1.5354
- val_sparse_categorical_accuracy: 0.7167 - lr: 0.0010

Epoch 2: LearningRateScheduler setting learning rate to 0.0009048374180359595.

Epoch 2/5

23/23 [=====] - ETA: 0s - loss: 6.6049 - sparse_categorical_accuracy: 0.7667

Epoch 2: val_loss improved from 1.53536 to 1.46090, saving model to /content/drive/MyDrive/model_1_cur.hdf5

23/23 [=====] - 16s 723ms/step - loss: 6.6049 - sparse_categorical_accuracy: 0.7667 - val_loss: 1.4609
- val_sparse_categorical_accuracy: 0.7167 - lr: 9.0484e-04

Epoch 3: LearningRateScheduler setting learning rate to 0.0008187307530779819.

Epoch 3/5

23/23 [=====] - ETA: 0s - loss: 3.4674 - sparse_categorical_accuracy: 0.8292

Epoch 3: val_loss improved from 1.46090 to 0.70322, saving model to /content/drive/MyDrive/model_1_cur.hdf5

23/23 [=====] - 16s 708ms/step - loss: 3.4674 - sparse_categorical_accuracy: 0.8292 - val_loss: 0.7032
- val_sparse_categorical_accuracy: 0.8944 - lr: 8.1873e-04

Epoch 4: LearningRateScheduler setting learning rate to 0.0007408182206817179.

Epoch 4/5

23/23 [=====] - ETA: 0s - loss: 4.0621 - sparse_categorical_accuracy: 0.8264

Epoch 4: val_loss did not improve from 0.70322

23/23 [=====] - 12s 519ms/step - loss: 4.0621 - sparse_categorical_accuracy: 0.8264 - val_loss: 0.8855
- val_sparse_categorical_accuracy: 0.8833 - lr: 7.4082e-04

Epoch 5: LearningRateScheduler setting learning rate to 0.0006703200460356394.

Epoch 5/5

23/23 [=====] - ETA: 0s - loss: 3.3155 - sparse_categorical_accuracy: 0.8653

Epoch 5: val_loss did not improve from 0.70322

23/23 [=====] - 12s 515ms/step - loss: 3.3155 - sparse_categorical_accuracy: 0.8653 - val_loss: 0.8769
- val_sparse_categorical_accuracy: 0.8889 - lr: 6.7032e-04

model_1 done

6/6 [=====] - 4s 133ms/step

weights done

Epoch 1: LearningRateScheduler setting learning rate to 0.001.

Epoch 1/5

23/23 [=====] - ETA: 0s - loss: 7.0102 - sparse_categorical_accuracy: 0.6458

Epoch 1: val_loss improved from inf to 2.09458, saving model to /content/drive/MyDrive/model_1_cur.hdf5

23/23 [=====] - 41s 918ms/step - loss: 7.0102 - sparse_categorical_accuracy: 0.6458 - val_loss: 2.0946
- val_sparse_categorical_accuracy: 0.7167 - lr: 0.0010

Epoch 2: LearningRateScheduler setting learning rate to 0.0009048374180359595.

Epoch 2/5

23/23 [=====] - ETA: 0s - loss: 5.1201 - sparse_categorical_accuracy: 0.7972

Epoch 2: val_loss improved from 2.09458 to 1.53172, saving model to /content/drive/MyDrive/model_1_cur.hdf5
23/23 [=====] - 16s 707ms/step - loss: 5.1201 - sparse_categorical_accuracy: 0.7972 - val_loss: 1.5317
- val_sparse_categorical_accuracy: 0.6389 - lr: 9.0484e-04

Epoch 3: LearningRateScheduler setting learning rate to 0.0008187307530779819.

Epoch 3/5

23/23 [=====] - ETA: 0s - loss: 3.7074 - sparse_categorical_accuracy: 0.8194
Epoch 3: val_loss improved from 1.53172 to 1.25236, saving model to /content/drive/MyDrive/model_1_cur.hdf5
23/23 [=====] - 17s 728ms/step - loss: 3.7074 - sparse_categorical_accuracy: 0.8194 - val_loss: 1.2524
- val_sparse_categorical_accuracy: 0.8389 - lr: 8.1873e-04

Epoch 4: LearningRateScheduler setting learning rate to 0.0007408182206817179.

Epoch 4/5

23/23 [=====] - ETA: 0s - loss: 3.6319 - sparse_categorical_accuracy: 0.8514
Epoch 4: val_loss improved from 1.25236 to 0.80076, saving model to /content/drive/MyDrive/model_1_cur.hdf5
23/23 [=====] - 17s 725ms/step - loss: 3.6319 - sparse_categorical_accuracy: 0.8514 - val_loss: 0.8008
- val_sparse_categorical_accuracy: 0.8278 - lr: 7.4082e-04

Epoch 5: LearningRateScheduler setting learning rate to 0.0006703200460356394.

Epoch 5/5

23/23 [=====] - ETA: 0s - loss: 2.0340 - sparse_categorical_accuracy: 0.9208
Epoch 5: val_loss did not improve from 0.80076
23/23 [=====] - 12s 514ms/step - loss: 2.0340 - sparse_categorical_accuracy: 0.9208 - val_loss: 0.8178
- val_sparse_categorical_accuracy: 0.8944 - lr: 6.7032e-04

model_1 done

6/6 [=====] - 4s 131ms/step

weights done

Epoch 1: LearningRateScheduler setting learning rate to 0.001.

Epoch 1/5

23/23 [=====] - ETA: 0s - loss: 6.7301 - sparse_categorical_accuracy: 0.6403
Epoch 1: val_loss improved from inf to 1.83547, saving model to /content/drive/MyDrive/model_1_cur.hdf5
23/23 [=====] - 42s 939ms/step - loss: 6.7301 - sparse_categorical_accuracy: 0.6403 - val_loss: 1.8355
- val_sparse_categorical_accuracy: 0.7056 - lr: 0.0010

Epoch 2: LearningRateScheduler setting learning rate to 0.0009048374180359595.

Epoch 2/5

23/23 [=====] - ETA: 0s - loss: 4.8875 - sparse_categorical_accuracy: 0.8264
Epoch 2: val_loss did not improve from 1.83547
23/23 [=====] - 12s 542ms/step - loss: 4.8875 - sparse_categorical_accuracy: 0.8264 - val_loss: 1.8554
- val_sparse_categorical_accuracy: 0.7889 - lr: 9.0484e-04

Epoch 3: LearningRateScheduler setting learning rate to 0.0008187307530779819.

Epoch 3/5

```
23/23 [=====] - ETA: 0s - loss: 3.6090 - sparse_categorical_accuracy: 0.8458
Epoch 3: val_loss improved from 1.83547 to 1.31485, saving model to /content/drive/MyDrive/model_1_cur.hdf5
23/23 [=====] - 17s 749ms/step - loss: 3.6090 - sparse_categorical_accuracy: 0.8458 - val_loss: 1.3149
- val_sparse_categorical_accuracy: 0.8000 - lr: 8.1873e-04

Epoch 4: LearningRateScheduler setting learning rate to 0.0007408182206817179.
Epoch 4/5
23/23 [=====] - ETA: 0s - loss: 2.7764 - sparse_categorical_accuracy: 0.8625
Epoch 4: val_loss improved from 1.31485 to 0.70417, saving model to /content/drive/MyDrive/model_1_cur.hdf5
23/23 [=====] - 16s 711ms/step - loss: 2.7764 - sparse_categorical_accuracy: 0.8625 - val_loss: 0.7042
- val_sparse_categorical_accuracy: 0.8722 - lr: 7.4082e-04

Epoch 5: LearningRateScheduler setting learning rate to 0.0006703200460356394.
Epoch 5/5
23/23 [=====] - ETA: 0s - loss: 1.6747 - sparse_categorical_accuracy: 0.9097
Epoch 5: val_loss improved from 0.70417 to 0.55556, saving model to /content/drive/MyDrive/model_1_cur.hdf5
23/23 [=====] - 16s 702ms/step - loss: 1.6747 - sparse_categorical_accuracy: 0.9097 - val_loss: 0.5556
- val_sparse_categorical_accuracy: 0.9000 - lr: 6.7032e-04
model_1 done
6/6 [=====] - 4s 134ms/step
weights done

Epoch 1: LearningRateScheduler setting learning rate to 0.001.
Epoch 1/5
23/23 [=====] - ETA: 0s - loss: 6.6773 - sparse_categorical_accuracy: 0.6542
Epoch 1: val_loss improved from inf to 2.14505, saving model to /content/drive/MyDrive/model_1_cur.hdf5
23/23 [=====] - 40s 885ms/step - loss: 6.6773 - sparse_categorical_accuracy: 0.6542 - val_loss: 2.1450
- val_sparse_categorical_accuracy: 0.6889 - lr: 0.0010

Epoch 2: LearningRateScheduler setting learning rate to 0.0009048374180359595.
Epoch 2/5
23/23 [=====] - ETA: 0s - loss: 5.5608 - sparse_categorical_accuracy: 0.7861
Epoch 2: val_loss did not improve from 2.14505
23/23 [=====] - 12s 515ms/step - loss: 5.5608 - sparse_categorical_accuracy: 0.7861 - val_loss: 2.2032
- val_sparse_categorical_accuracy: 0.7333 - lr: 9.0484e-04


Epoch 3: LearningRateScheduler setting learning rate to 0.0008187307530779819.
Epoch 3/5
23/23 [=====] - ETA: 0s - loss: 4.3022 - sparse_categorical_accuracy: 0.8028
Epoch 3: val_loss improved from 2.14505 to 0.59682, saving model to /content/drive/MyDrive/model_1_cur.hdf5
23/23 [=====] - 16s 722ms/step - loss: 4.3022 - sparse_categorical_accuracy: 0.8028 - val_loss: 0.5968
- val_sparse_categorical_accuracy: 0.8278 - lr: 8.1873e-04

Epoch 4: LearningRateScheduler setting learning rate to 0.0007408182206817179.
```



```
Epoch 4/5
23/23 [=====] - ETA: 0s - loss: 2.3015 - sparse_categorical_accuracy: 0.8556
Epoch 4: val_loss improved from 0.59682 to 0.58581, saving model to /content/drive/MyDrive/model_1_cur.hdf5
23/23 [=====] - 17s 725ms/step - loss: 2.3015 - sparse_categorical_accuracy: 0.8556 - val_loss: 0.5858
- val_sparse_categorical_accuracy: 0.9056 - lr: 7.4082e-04

Epoch 5: LearningRateScheduler setting learning rate to 0.0006703200460356394.
Epoch 5/5
23/23 [=====] - ETA: 0s - loss: 2.0464 - sparse_categorical_accuracy: 0.8986
Epoch 5: val_loss did not improve from 0.58581
23/23 [=====] - 12s 542ms/step - loss: 2.0464 - sparse_categorical_accuracy: 0.8986 - val_loss: 0.7394
- val_sparse_categorical_accuracy: 0.8722 - lr: 6.7032e-04
model_1 done
6/6 [=====] - 4s 132ms/step
```



Использование аугментации данных

В нейронной сети используется слой аугментации данных (поворот вертикальный/горизонтальный, поворот на угол, увеличение масштаба изображения).

Выбраны параметры, не слишком сильно увеличивающие изображение (чтобы не терять в качестве). Zoom выбран исключительно отрицательным чтобы не дополнять изображение паддингом (т.к. тогда создаются зеркальные края, не всегда отражающие реальную структуру ткани).


```
In [ ]: import keras
        from keras import layers
```

```
import matplotlib.pyplot as plt

plt.figure(figsize=(3, 3))
images, lbl = d_train.random_batch_with_labels(1)
print(images.shape)
ax = plt.subplot(1, 1, 1)
plt.imshow(images[0].astype("uint8"))
plt.axis("off")
```










```
(1, 224, 224, 3)
(-0.5, 223.5, 223.5, -0.5)
```

Out[]:



```
In [ ]: data_augmentation = tf.keras.Sequential(
    [
        tf.keras.layers.RandomFlip(),
        tf.keras.layers.RandomRotation(0.2),
        tf.keras.layers.RandomZoom((-0.2, 0)),
    ]
)
data_augmentation.compile(optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics=['accuracy'])
```

```
In [ ]: plt.figure(figsize=(10, 10))
for i in range(9):
    augmented_images = data_augmentation(images, 0.1)
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(augmented_images[0].numpy().astype("uint8"))
    plt.axis("off")
```



In []:

Пример тестирования модели на полном наборе данных:

```
In [ ]: # evaluating model on full test dataset (may take time)
if TEST_ON_LARGE_DATASET:
```

```
pred_2 = model.test_on_dataset(d_test)
Metrics.print_all(d_test.labels, pred_2, 'test')
```

Результат работы пайплайна обучения и тестирования выше тоже будет оцениваться. Поэтому не забудьте присылать на проверку ноутбук с выполненными ячейками кода с демонстрациями метрик обучения, графиками и т.п. В этом пайплайне Вам необходимо продемонстрировать работу всех реализованных дополнений, улучшений и т.п.

Настоятельно рекомендуется после получения пайплайна с полными результатами обучения экспортировать ноутбук в pdf (файл -> печать) и прислать этот pdf вместе с самим ноутбуком.

Тестирование модели на других наборах данных

Ваша модель должна поддерживать тестирование на других наборах данных. Для удобства, Вам предоставляется набор данных test_tiny, который представляет собой малую часть (2% изображений) набора test. Ниже приведен фрагмент кода, который будет осуществлять тестирование для оценивания Вашей модели на дополнительных тестовых наборах данных.

Прежде чем отсылать задание на проверку, убедитесь в работоспособности фрагмента кода ниже.

```
In [ ]: final_model = Model()
final_model.load('best')
d_test_tiny = Dataset('test_tiny')
pred = final_model.test_on_dataset(d_test_tiny)
Metrics.print_all(d_test_tiny.labels, pred, 'test_tiny')
```

loading started

Downloading...

From: <https://drive.google.com/uc?id=1-8StHduxeVrtZAE1-5AQIPT0YQj11Rso>

To: /content/best.hdf5

100%|██████████| 251M/251M [00:00<00:00, 293MB/s]

loading done

Downloading...

From: <https://drive.google.com/uc?export=download&confirm=pbef&id=1JcDlKmFQ-ohdFVMQDINGY2EauuSGSVgv>

To: /content/test_tiny.npz

100%|██████████| 10.6M/10.6M [00:00<00:00, 16.6MB/s]

```
Loading dataset test_tiny from npz.  
Done. Dataset test_tiny consists of 90 images.  
3/3 [=====] - 4s 484ms/step  
metrics for test_tiny:  
    accuracy 0.9889:  
    balanced accuracy 0.9889:
```

Отмонтировать Google Drive.

```
In [ ]: drive.flush_and_unmount()
```

Дополнительные "полезности"

Ниже приведены примеры использования различных функций и библиотек, которые могут быть полезны при выполнении данного практического задания.

Измерение времени работы кода

Измерять время работы какой-либо функции можно легко и непринужденно при помощи функции `timeit` из соответствующего модуля:

```
In [ ]: import timeit  
  
def factorial(n):  
    res = 1  
    for i in range(1, n + 1):  
        res *= i  
    return res  
  
def f():  
    return factorial(n=1000)  
  
n_runs = 128  
print(f'Function f is calculated {n_runs} times in {timeit.timeit(f, number=n_runs)}s.')
```

Scikit-learn

Для использования "классических" алгоритмов машинного обучения рекомендуется использовать библиотеку scikit-learn (<https://scikit-learn.org/stable/>). Пример классификации изображений цифр из набора данных MNIST при помощи классификатора SVM:

```
In [ ]: # Standard scientific Python imports
import matplotlib.pyplot as plt

# Import datasets, classifiers and performance metrics
from sklearn import datasets, svm, metrics
from sklearn.model_selection import train_test_split

# The digits dataset
digits = datasets.load_digits()

# The data that we are interested in is made of 8x8 images of digits, let's
# have a look at the first 4 images, stored in the `images` attribute of the
# dataset. If we were working from image files, we could load them using
# matplotlib.pyplot.imread. Note that each image must have the same size. For these
# images, we know which digit they represent: it is given in the 'target' of
# the dataset.
_, axes = plt.subplots(2, 4)
images_and_labels = list(zip(digits.images, digits.target))
for ax, (image, label) in zip(axes[0, :], images_and_labels[:4]):
    ax.set_axis_off()
    ax.imshow(image, cmap=plt.cm.gray_r, interpolation='nearest')
    ax.set_title('Training: %i' % label)

# To apply a classifier on this data, we need to flatten the image, to
# turn the data in a (samples, feature) matrix:
n_samples = len(digits.images)
data = digits.images.reshape((n_samples, -1))

# Create a classifier: a support vector classifier
classifier = svm.SVC(gamma=0.001)

# Split data into train and test subsets
X_train, X_test, y_train, y_test = train_test_split(
    data, digits.target, test_size=0.5, shuffle=False)

# We learn the digits on the first half of the digits
classifier.fit(X_train, y_train)
```

```
# Now predict the value of the digit on the second half:
predicted = classifier.predict(X_test)

images_and_predictions = list(zip(digits.images[n_samples // 2:], predicted))
for ax, (image, prediction) in zip(axes[1, :], images_and_predictions[:4]):
    ax.set_axis_off()
    ax.imshow(image, cmap=plt.cm.gray_r, interpolation='nearest')
    ax.set_title('Prediction: %i' % prediction)

print("Classification report for classifier %s:\n%s\n"
      % (classifier, metrics.classification_report(y_test, predicted)))
disp = metrics.plot_confusion_matrix(classifier, X_test, y_test)
disp.figure_.suptitle("Confusion Matrix")
print("Confusion matrix:\n%s" % disp.confusion_matrix)

plt.show()
```

Scikit-image

Реализовывать различные операции для работы с изображениями можно как самостоятельно, работая с массивами numpy, так и используя специализированные библиотеки, например, scikit-image (<https://scikit-image.org/>). Ниже приведен пример использования Canny edge detector.

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
from scipy import ndimage as ndi
import cv2
from skimage import feature

# Generate noisy image of a square
im = d_train.images[np.random.randint(d_train.images.shape[0])][:,:,0]

# Compute the Canny filter for two values of sigma
edges1 = feature.canny(im)
edges2 = feature.canny(im, sigma=2)
ridge_filter = cv2.ximgproc.RidgeDetectionFilter_create()
edges3 = ridge_filter.getRidgeFilteredImage(im)
# display results
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(nrows=2, ncols=2, figsize=(8, 3),
```

```
sharex=True, sharey=True)

fig.set_size_inches(10, 10)

ax1.imshow(im, cmap=plt.cm.gray)
ax1.axis('off')
ax1.set_title('noisy image', fontsize=20)

ax2.imshow(edges1, cmap=plt.cm.gray)
ax2.axis('off')
ax2.set_title(r'Canny filter,  $\sigma=1$ ', fontsize=20)

ax3.imshow(edges2, cmap=plt.cm.gray)
ax3.axis('off')
ax3.set_title(r'Canny filter,  $\sigma=2$ ', fontsize=20)

ax4.imshow(edges3, cmap=plt.cm.gray)
ax4.axis('off')
ax4.set_title(r'Canny filter,  $\sigma=3$ ', fontsize=20)

fig.tight_layout()

plt.show()
```

```
In [ ]: plt.figure(figsize=(3, 3))
images, lbl = d_train.random_batch_with_labels(1)
print(images.shape)
ax = plt.subplot(1, 1, 1)
plt.imshow(images[0].astype("uint8"))
plt.axis("off")
```

Tensorflow 2

Для создания и обучения нейросетевых моделей можно использовать фреймворк глубокого обучения Tensorflow 2. Ниже приведен пример простейшей нейронной сети, использующейся для классификации изображений из набора данных MNIST.

```
In [ ]: # Install TensorFlow

import tensorflow as tf

mnist = tf.keras.datasets.mnist
```



```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)

model.evaluate(x_test, y_test, verbose=2)
```

Для эффективной работы с моделями глубокого обучения убедитесь в том, что в текущей среде Google Colab используется аппаратный ускоритель GPU или TPU. Для смены среды выберите "среда выполнения" -> "сменить среду выполнения".

Большое количество tutorиалов и примеров с кодом на Tensorflow 2 можно найти на официальном сайте <https://www.tensorflow.org/tutorials?hl=ru>.

Также, Вам может понадобиться написать собственный генератор данных для Tensorflow 2. Скорее всего он будет достаточно простым, и его легко можно будет реализовать, используя официальную документацию TensorFlow 2. Но, на всякий случай (если не удалось сразу разобраться или хочется вникнуть в тему более глубоко), можете посмотреть следующий отличный tutorиал: <https://stanford.edu/~shervine/blog/keras-how-to-generate-data-on-the-fly>.

Numba

В некоторых ситуациях, при ручных реализациях графовых алгоритмов, выполнение многократных вложенных циклов for в python можно существенно ускорить, используя JIT-компилятор Numba (<https://numba.pydata.org/>). Примеры использования Numba в Google Colab можно найти тут:

1. https://colab.research.google.com/github/cbnet/maldives/blob/master/numba/numba_cuda.ipynb
2. https://colab.research.google.com/github/evaneschneider/parallel-programming/blob/master/COMPASS_gpu_intro.ipynb

Пожалуйста, если Вы решили использовать Numba для решения этого практического задания, еще раз подумайте, нужно ли это Вам, и есть ли возможность реализовать требуемую функциональность иным способом. Используйте Numba только при реальной необходимости.

Работа с zip архивами в Google Drive

Запаковка и распаковка zip архивов может пригодиться при сохранении и загрузки Вашей модели. Ниже приведен фрагмент кода, иллюстрирующий помещение нескольких файлов в zip архив с последующим чтением файлов из него. Все действия с директориями, файлами и архивами должны осущетвляться с примонтированным Google Drive.

Создадим 2 изображения, поместим их в директорию tmp внутри PROJECT_DIR, запакуем директорию tmp в архив tmp.zip.

```
In [ ]: PROJECT_DIR = "/dev/prak_nn_1/"
arr1 = np.random.rand(100, 100, 3) * 255
arr2 = np.random.rand(100, 100, 3) * 255

img1 = Image.fromarray(arr1.astype('uint8'))
img2 = Image.fromarray(arr2.astype('uint8'))

p = "/content/drive/MyDrive/" + PROJECT_DIR

if not (Path(p) / 'tmp').exists():
    (Path(p) / 'tmp').mkdir()

img1.save(str(Path(p) / 'tmp' / 'img1.png'))
img2.save(str(Path(p) / 'tmp' / 'img2.png'))

%cd $p
!zip -r "tmp.zip" "tmp"
```

Распакуем архив tmp.zip в директорию tmp2 в PROJECT_DIR. Теперь внутри директории tmp2 содержится директория tmp, внутри которой находятся 2 изображения.

```
In [ ]: p = "/content/drive/MyDrive/" + PROJECT_DIR
%cd $p
!unzip -uq "tmp.zip" -d "tmp2"
```